Shortlist AIO — Business Plan

An AI answer-visibility firm for B2B SaaS companies.

Trading name: **Shortlist AIO**. Client-facing everywhere — site, contracts, reports, invoices. **Cold-outbound sending domains stay topically neutral** and never carry the brand; see `../ops/EMAIL-INFRASTRUCTURE.md` §1 for why. Operated end-to-end from a computer. Claude does the work; the human principal is the legal person, signatory, and account holder.

Version	1.1 — 25 August 2026
Starting capital	$1,000, then reinvested profit only
Objective	Maximize recognized revenue over 12 months
Base-case target	~$387,000 Year-1 revenue; ~$70,900 exit MRR (~$851k ARR run-rate)
Human time required	3–6 hours/week

---

1. The one-sentence business

We make B2B software companies the vendor that ChatGPT, Perplexity, Gemini, and Google AI Overviews name when a buyer asks *"what's the best tool for X?"* — sold as a fixed-price monthly program, priced against human agencies, delivered at a fraction of their cost.

---

2. Why this business, and not another one

I evaluated six candidate models against four filters: (1) demand that already has budget, (2) distribution I can execute without ad spend or an existing audience, (3) delivery that is pure computer-based knowledge work, (4) revenue inside 12 months, not 36.

2.1 The demand shock is real, recent, and measurable

The buying behavior of B2B software purchasers moved faster than vendor marketing did:

- **72%** of B2B software buyers now use ChatGPT as part of vendor evaluation; **51%** start research with an AI chatbot more often than with Google.
- **69%** chose a different vendor than originally planned because of AI chatbot guidance; **33%** bought from a vendor they had not previously heard of.
- The average vendor shortlist has compressed to **2.5 names**, down from 3.2.
- And the supply side has not caught up: **51% of B2B tech brands have zero citations** across ChatGPT, Perplexity, and Gemini.

That last number is the whole business. Half the addressable market has a problem that is urgent, expensive, and — critically — **provable in ten minutes with a script**. I can show a prospect the transcript where a named competitor is recommended and they are not.

Meanwhile the budget that used to fund this outcome is stranded. Where an AI Overview appears, top-ranking pages lost **58%** of click-through; **60%** of Google searches now end with no click at all. Marketing leaders have line items that no longer produce pipeline and are actively looking for somewhere to move them.

2.2 The money is already moving

Benchmark	2026 market rate
GEO / AI-visibility services market	~$1.5–2.0B in 2026, projected $4B+ by 2027
SMB GEO retainer	$1,500–$5,000/mo
Mid-market GEO retainer	$2,000–$8,000/mo
One-time AI visibility audit	$500–$2,000
Digital PR retainer	$3,000–$30,000/mo
Single earned editorial placement	$750–$2,000
Original research campaign	$5,000–$15,000

We are not creating a category or educating a market. We are entering a category with established, published price points and undercutting the labor cost by ~95%.

2.3 The work that actually moves the metric is exactly what an AI does best

This matters, because a service that doesn't work is a refund pipeline. The published correlation data says what drives citations:

- **Brand web mentions correlate 0.664** with AI citation rate — roughly **3× stronger than backlinks** (0.218).
- **~75% of the citation equation lives on external domains**, not the client's own site.
- Domains with heavy Reddit brand presence average **3.9× more ChatGPT citations**.
- Platforms diverge sharply — only **11% of domains** are cited by both ChatGPT and Perplexity; ChatGPT cites brands 0.59% of the time vs Perplexity's 13.05%. Per-engine strategy is required, which is precisely the kind of tedious, high-volume differentiation human agencies skip.

So the deliverable is: research, structured content at volume, entity and data hygiene across dozens of third-party surfaces, original data studies, and per-engine tuning. That is labor. Human agencies charge $5–15k/month for it because it *is* labor. My marginal cost to produce it is roughly $150/month per client.

That gap — a $2,500–$5,000 price anchored to human labor against a ~$150 cost to serve — is the entire thesis.

2.4 The sales motion is the part most AI businesses get wrong, and this one is native to me

Distribution, not production, is the binding constraint for any AI business in 2026. Everyone can produce. Here the outreach *is* the product demo:

> *"We asked ChatGPT, Perplexity, and Gemini the 40 questions your buyers ask before choosing a [category] tool. You were named in 2 of 120 answers. [Competitor] was named in 81. Here's the transcript."*

That email is bespoke, verifiable, and impossible to argue with. Producing it requires reading a prospect's whole web presence and running 120 API calls — trivial for me, uneconomic for a human SDR. Benchmark cold email replies at 3.1–3.7%; a report-backed email should clear that, and we model 4%.

And B2B SaaS buyers **transact asynchronously** — email, a Stripe link, a shared Slack channel. No phone call is required to close a $1,500 diagnostic. For a business where the operator cannot take calls, this is decisive.

2.5 What I rejected, and why

Rejected	Reason
Affiliate / SEO content sites	Structurally broken. 71% of affiliate sites lost rankings in the March 2026 core update; personal finance -99%, fashion -95%. Building on a collapsing channel.
Micro-SaaS product	70% never pass $1k/mo; 40% never reach $1k MRR at all; median time to $1M ARR is 2 years 9 months. Wrong instrument for a 12-month revenue goal, and the cold-start distribution problem is my single greatest weakness.
An AI-visibility tracking tool	Already commoditized. $300M+ raised in the category since mid-2025; Profound at a $1B valuation on $155M; usable tools start at $25/mo. Never compete with venture-funded commodity infrastructure — **consume it instead**.
Shopify apps	17,600 apps, median under $1,000/mo, 5–12% *monthly* churn. Crowded and leaky.
ADA / accessibility remediation	Genuine regulatory demand, but DOJ pushed Title II to April 2027/2028, selling "compliance" carries direct liability, and DOJ explicitly cited generative AI's inability to automate remediation at scale. Keep as a Phase 3 adjacency, not the beachhead.
Upwork / Fiverr freelancing	Price-taking, platform-dependent, capped, and requires misrepresenting who does the work.

---

3. Customer and offer

3.1 Ideal customer profile

A B2B SaaS company that **already buys marketing** and is **demonstrably invisible in AI answers**:

- Series A–C, 20–200 employees, roughly $3M–$50M ARR
- Has a G2/Capterra presence and an active content program (proves budget and buy-in exist)
- Sells into a category with a real "best X software" query pattern
- Our scan shows <10% share of voice while a direct competitor exceeds 40%

Beachhead categories (chosen for vendor density, budget, low agency saturation, and — see §8 — insulation from the goods trade war): **legal tech** and **compliance/regtech**, with **cybersecurity** as the third. These serve buyers whose own demand is non-discretionary: regulatory deadlines and legal exposure do not soften when the economy does.

Revised August 2026. Earlier drafts named construction tech and field service as the beachhead. That is now the wrong call. Those vendors sell to contractors buying steel, copper, and building materials at 50% tariffs — margin compression flows straight through to software churn and frozen marketing budgets. HR/recruiting tech is demoted to secondary for a milder version of the same problem: it is levered to hiring, which is the first thing to slow. Expand once each category's research assets and case studies compound.

3.2 Offer ladder

Tier	Price	What it is	Purpose
AI Visibility Report	Free	Automated 120-query scan across 4 engines, share-of-voice vs 3 competitors, transcripts	The cold-email hook. Costs ~$0.40 to produce.
Diagnostic Sprint	**$1,500** one-time	2 weeks: 200-prompt tracked set, competitor citation teardown, source-gap map, 90-day roadmap	De-risks the buyer, produces cash immediately, converts to retainer at ~50%
Visibility Program	**$2,500/mo**	Tracked prompt set; 8–12 answer-engine content assets/mo; schema + entity hygiene; third-party profile coverage; monthly report	Core recurring product
Authority Program	**$5,000/mo**	Everything above + one original research study per quarter + digital PR pitching + comparison-page inclusion outreach	Attacks the 75% of citations that live off-site
Category Ownership	**$8,500/mo**	Multi-category/multi-geo, proprietary index or dataset, monthly study, quarterly strategy session	Ceiling-raiser for the best accounts

Contracts: month-to-month after a 3-month initial term, 30 days' notice. **Deliverable-based SLA — never an outcome guarantee.** We commit to activities and reporting, not to a citation rate we do not control.

3.3 Unit economics (Visibility Program, $2,500/mo)

Line	Monthly
Revenue	$2,500
LLM API — content production (Claude Sonnet 5 at $3/$15 per MTok, Opus 5 at $5/$25 for high-value assets)	~$110
Search / citation-tracking API (Perplexity Search API at $5 per 1,000 requests; ~960 tracked queries/mo)	~$10
Allocated shared tooling and infrastructure	~$30
Gross margin	**~$2,350 (94%)**

Customer acquisition cost is dominated by outbound infrastructure at roughly $150/month total, spread across every deal — well under $200 per client acquired at steady state. The dominant constraint is not money. It is delivery quality at volume, which Section 6 addresses.

---

4. Financial plan

Full month-by-month build, funnel assumptions, and scenario mechanics are in `FINANCIAL-MODEL.md`. Summary:

Scenario	New retainers/mo (steady)	Exit MRR	Year-1 revenue
Conservative	1.5	~$29,000	~$159,000
Base	**3.0**	**~$70,900**	**~$387,000**
Aggressive	4.0 @ higher tier mix (capacity-capped)	~$115,400	~$698,000

Base case assumes 3,000 emails/month at steady state, a 4.0% reply rate, 25% of replies positive, 22% of conversations closing a paid diagnostic, 50% diagnostic-to-retainer conversion, a $3,200 blended retainer, and **5% monthly churn**. It ends the year with roughly 22 active retainer clients.

Opening budget — the $1,000

Item	Cost
Entity formation	$0 — Manitoba corporation already held
US-domiciled USD account setup + cross-border accounting reserve	$150
Primary domain + 8 sending domains	$110
10 sending mailboxes, 3 months prepaid	$75
Smartlead, 3 months ($39/mo, unlimited mailboxes)	$117
LLM API credits	$200
Search / citation API credits	$50
Lead enrichment credits	$100
Brand assets, hosting, miscellaneous	$50
Unallocated reserve	$148
Total	**$1,000**

Website, tracking dashboard, outreach tooling, report generator, and every contract draft are built by me at no cash cost. Hosting is on a free tier. Stripe takes 2.9% + $0.30 with nothing upfront.

Reinvestment rule: through month 6, 100% of profit goes to sending capacity and API credits, in that order — outreach volume is the binding growth constraint. From month 7, cap outbound spend and route surplus into (a) original research studies that generate inbound and (b) a modest reserve against refunds. Take no distributions before month 9.

---

5. Go-to-market

Primary channel — report-led cold outbound. Build a target list of ~4,000 qualified companies by scanning public sources (G2 category pages, funding announcements, job postings for marketing roles). Score each one by running the free scan. Contact only companies whose scan shows a real, specific gap — this doubles as list hygiene and keeps spam complaints structurally low, because every email is genuinely relevant to its recipient.

Volume ramps as deliverability allows: 200/day maximum per new domain, 30-day ramp, targeting ~150/day steady state across ~10 mailboxes.

Secondary — publish the aggregate data. Every scan we run feeds a proprietary dataset: who gets cited, for what queries, on which engine, in which category. Quarterly we publish *The B2B AI Citation Index* from it. This does three things at once: it earns the citations and mentions that prove our own method works, it generates inbound, and it is the single most defensible asset the business builds. **The marketing channel is a live demonstration of the product.**

Tertiary — proof assets. Three free pilot clients in months 1–2, chosen for logo quality, in exchange for a case study and a testimonial. This is the fastest route past the no-track-record problem, and it is honest: real work, real results, disclosed relationship.

---

6. Operating model — how this actually runs

I do not run continuously. I run when invoked. Every recurring process must therefore be **code, not a Claude session**:

- Nightly scheduled jobs (GitHub Actions, free tier) run citation tracking for every client and every prospect, write results to a database, and flag movement.
- Report generation is templated and automated; I write the analysis and recommendations, not the assembly.
- The outreach queue is populated automatically; I write the copy.

Who	Does what	Time
Claude	All research, content, code, audits, reports, outreach copy, contract drafts, client email drafts	Invoked in batches
Principal (human)	Signs contracts, holds bank/Stripe/domain accounts, approves outbound before send, takes occasional calls, runs delivery sessions, final quality gate	3–6 hrs/week

The binding constraint is delivery quality at scale, not lead generation. At 22 clients the program produces 200–260 content assets per month. The correct response to demand beyond that is **to raise prices, not to add clients** — a 20% price increase and a 10% client reduction is strictly better for both margin and quality. Plan to raise prices at month 6 and month 10.

---

7. Legal, compliance, and ethics

Detail and the documents for counsel are in `COMPLIANCE-POLICY.md`. The essentials:

Entity — a Manitoba share corporation. The principal already holds one, and it is the right vehicle, not merely an acceptable one. Selling services into the US from a Canadian company is routine and does not meaningfully affect market access at this deal size. Three consequences matter:

- **Trade war exposure: none, and the currency helps.** Tariffs are an instrument for goods; they do not apply to cross-border services or intangibles, and the September 2026 Canadian retaliation targets steel, dairy, agricultural equipment, appliances, paper, and electronics. Nothing we sell appears on any tariff schedule. Meanwhile a weak loonie (USD/CAD ~1.40) is a tailwind for a Canadian company invoicing in USD. Canada's Digital Services Tax — the one measure that would have reached a firm like this — was rescinded in 2025 and repealed retroactively, with refunds issued.
- **Tax.** Manitoba levies **0% provincial tax** on the first $500,000 of active business income for a Canadian-Controlled Private Corporation; with the 9% federal small-business rate that is a **9% combined rate**, the lowest in Canada alongside Yukon. The base case's ~$353,000 pre-tax profit sits entirely inside that band. This is a deferral advantage rather than an absolute one — personal tax applies when profits are distributed — but the plan reinvests through month 6 and holds reserves after, so the deferral is real money. *Condition:* CCPC status requires Canadian control. Without it the general rate applies (12% Manitoba + 15% federal = 27%).
- **No US federal income tax, no US withholding.** Under Articles V and VII of the Canada–US treaty, business profits are taxable in the US only through a US permanent establishment. Services performed from Canada for US clients create none. Separately, services performed outside the US are foreign-source income and are not subject to the 30% withholding — clients receive a **Form W-8BEN-E** and withhold nothing. The treaty position is not automatic: it requires a **timely Form 1120-F protective return with Form 8833 attached**, filed annually.
- **GST/HST works in our favour.** Advisory, consulting, and professional services supplied to non-resident, non-registered persons are **zero-rated** (Schedule VI, Part V, s. 23). We charge US clients 0%, and because zero-rated is not the same as exempt, we still claim input tax credits on Canadian expenses. Registration is required above $30,000 CAD in four consecutive quarters and is worth doing voluntarily before then to recover ITCs.

Detail, and the questions for counsel and an accountant, are in `COMPLIANCE-POLICY.md`.

Hard ethical lines — non-negotiable, and also the correct business decision. The FTC's Rule on Consumer Reviews and Testimonials (effective 21 October 2024) carries civil penalties up to roughly **$52,000 per violation** and reaches conduct a company "knew or should have known" about. Accordingly:

- We never write, buy, sell, or incentivize reviews or testimonials — including AI-generated ones.
- We never post as a client or operate sockpuppet accounts. Community participation is done by named client employees with disclosed affiliation.
- We never build company-controlled sites presenting themselves as independent comparison or review sources.

- We disclose to every client, in the MSA, that delivery is AI-assisted.

This is not merely compliance posture. Astroturfing is the obvious shortcut in this category, it is the reason the category will eventually attract enforcement, and a firm that publishes its methodology and refuses to fake anything is differentiated on exactly the axis clients will start caring about.

Outbound. CAN-SPAM: accurate headers, honest subject lines, physical postal address, opt-outs honored within 10 business days. GDPR: B2B legitimate interest with a documented Legitimate Interest Assessment, disclosed data source, one-click opt-out. Sender rules: SPF, DKIM, and aligned DMARC; spam complaints held under 0.10%; 14-day minimum domain age; 30-day volume ramp.

Measurement. Citation tracking uses official paid APIs, never scraping against platform terms.

---

8. Risks and kill criteria

Risk	Severity	Mitigation
Cold outbound underperforms without a track record	**High**	Three free pilots for case studies in months 1–2; the report hook is the strongest available differentiator; kill criteria below
GEO commoditizes; tools automate the on-site layer	High	Move up the ladder into off-site earned mentions, original research, and PR — the 75% that cannot be automated by a dashboard
LLM platforms change citation mechanics	Medium	Contracts commit to activities, not outcomes; the tracking dataset detects shifts before clients do
Delivery quality degrades past ~20 clients	Medium	Cap client count; raise prices at months 6 and 10; automate assembly, never analysis
Client churn after 90 days without visible lift	Medium	Set expectations at 90 days in writing; report leading indicators (mention volume, source coverage) monthly, not just citations
Deliverability collapse	Medium	Ramp discipline, scan-qualified lists only, complaint monitoring, spare domains warmed in reserve
US–Canada trade war depresses SaaS marketing budgets	Medium	No direct exposure — tariffs apply to goods, not services — but a slowing economy is a real threat to the base case. Partly hedged by the nature of the pitch: we sell *reallocation* of an existing budget that is visibly underperforming, not a new line item. Reallocation pitches hold up better in a downturn than incremental-spend pitches. Beachhead verticals chosen for goods-exposure insulation (§3.1).
US buyer reluctance toward a Canadian supplier	Low	Evidence says B2B relationships are holding; a service carries no tariff cost to the buyer, so our nationality is commercially irrelevant to them. Neutral `.com` domain, USD pricing, and outreach that leads with their data rather than our location. Disclosed plainly in the footer, MSA, and privacy policy — not led with.

Risk	Severity	Mitigation
Non-tariff measures restrict access to US AI APIs	Low	Tail risk only: services cannot be tariffed, and negotiations are reportedly moving toward *more* cross-border data freedom. The engine layer is already pluggable adapters — losing one provider costs an engine, not the business.
Single-operator concentration	Low	Everything in version control and documented; no undocumented process

Kill criteria — decide, do not drift. If at the end of **month 4** the business has fewer than 2 paying retainers and under $8,000 cumulative revenue, stop selling and reassess. The most likely correction is a vertical change, not a business-model change: the same machine repoints at a different category in about two weeks.

---

9. What happens after year one

The two-year path is not "more retainers." It is:

1. **Productize the diagnostic** into a self-serve $299 report — same automated pipeline, no human in the loop, pure margin, and a top-of-funnel machine.
2. **Publish the Citation Index** as a durable, cited data asset that makes the firm the source rather than the vendor.
3. **Move up-market.** The same program sold to a $100M+ ARR company is a $15,000–$25,000/month engagement, and there are fewer than 30 firms in the world who will have a two-year citation dataset to sell alongside it.

Sources

Market and benchmark figures cited above are drawn from the research recorded in `RESEARCH-NOTES.md`, which lists each source URL against the claim it supports.