Shortlist AIO — Business Plan

An AI answer-visibility firm for B2B SaaS companies.

Trading name: **Shortlist AIO**. Client-facing everywhere — site, contracts, reports, invoices. **Cold-outbound sending domains stay topically neutral** and never carry the brand; see `../04-Setup-Guides/EMAIL-INFRASTRUCTURE.md` §1 for why. Operated end-to-end from a computer. Claude does the work; the human principal is the legal person, signatory, and account holder.

Version	1.1 — 25 August 2026
Starting capital	$1,000, then reinvested profit only
Objective	Maximize recognized revenue over 12 months
Base-case target	~$387,000 Year-1 revenue; ~$70,900 exit MRR (~$851k ARR run-rate)
Human time required	3–6 hours/week

---

1. The one-sentence business

We make B2B software companies the vendor that ChatGPT, Perplexity, Gemini, and Google AI Overviews name when a buyer asks *"what's the best tool for X?"* — sold as a fixed-price monthly program, priced against human agencies, delivered at a fraction of their cost.

---

2. Why this business, and not another one

I evaluated six candidate models against four filters: (1) demand that already has budget, (2) distribution I can execute without ad spend or an existing audience, (3) delivery that is pure computer-based knowledge work, (4) revenue inside 12 months, not 36.

2.1 The demand shock is real, recent, and measurable

The buying behavior of B2B software purchasers moved faster than vendor marketing did:

- **72%** of B2B software buyers now use ChatGPT as part of vendor evaluation; **51%** start research with an AI chatbot more often than with Google.
- **69%** chose a different vendor than originally planned because of AI chatbot guidance; **33%** bought from a vendor they had not previously heard of.
- The average vendor shortlist has compressed to **2.5 names**, down from 3.2.
- And the supply side has not caught up: **51% of B2B tech brands have zero citations** across ChatGPT, Perplexity, and Gemini.

That last number is the whole business. Half the addressable market has a problem that is urgent, expensive, and — critically — **provable in ten minutes with a script**. I can show a prospect the transcript where a named competitor is recommended and they are not.

Meanwhile the budget that used to fund this outcome is stranded. Where an AI Overview appears, top-ranking pages lost **58%** of click-through; **60%** of Google searches now end with no click at all. Marketing leaders have line items that no longer produce pipeline and are actively looking for somewhere to move them.

2.2 The money is already moving

Benchmark	2026 market rate
GEO / AI-visibility services market	~$1.5–2.0B in 2026, projected $4B+ by 2027
SMB GEO retainer	$1,500–$5,000/mo
Mid-market GEO retainer	$2,000–$8,000/mo
One-time AI visibility audit	$500–$2,000
Digital PR retainer	$3,000–$30,000/mo
Single earned editorial placement	$750–$2,000
Original research campaign	$5,000–$15,000

We are not creating a category or educating a market. We are entering a category with established, published price points and undercutting the labor cost by ~95%.

2.3 The work that actually moves the metric is exactly what an AI does best

This matters, because a service that doesn't work is a refund pipeline. The published correlation data says what drives citations:

- **Brand web mentions correlate 0.664** with AI citation rate — roughly **3× stronger than backlinks** (0.218).
- **~75% of the citation equation lives on external domains**, not the client's own site.
- Domains with heavy Reddit brand presence average **3.9× more ChatGPT citations**.
- Platforms diverge sharply — only **11% of domains** are cited by both ChatGPT and Perplexity; ChatGPT cites brands 0.59% of the time vs Perplexity's 13.05%. Per-engine strategy is required, which is precisely the kind of tedious, high-volume differentiation human agencies skip.

So the deliverable is: research, structured content at volume, entity and data hygiene across dozens of third-party surfaces, original data studies, and per-engine tuning. That is labor. Human agencies charge $5–15k/month for it because it *is* labor. My marginal cost to produce it is roughly $150/month per client.

That gap — a $2,500–$5,000 price anchored to human labor against a ~$150 cost to serve — is the entire thesis.

2.4 The sales motion is the part most AI businesses get wrong, and this one is native to me

Distribution, not production, is the binding constraint for any AI business in 2026. Everyone can produce. Here the outreach *is* the product demo:

> *"We asked ChatGPT, Perplexity, and Gemini the 40 questions your buyers ask before choosing a [category] tool. You were named in 2 of 120 answers. [Competitor] was named in 81. Here's the transcript."*

That email is bespoke, verifiable, and impossible to argue with. Producing it requires reading a prospect's whole web presence and running 120 API calls — trivial for me, uneconomic for a human SDR. Benchmark cold email replies at 3.1–3.7%; a report-backed email should clear that, and we model 4%.

And B2B SaaS buyers **transact asynchronously** — email, a Stripe link, a shared Slack channel. No phone call is required to close a $1,500 diagnostic. For a business where the operator cannot take calls, this is decisive.

2.5 What I rejected, and why

Rejected	Reason
Affiliate / SEO content sites	Structurally broken. 71% of affiliate sites lost rankings in the March 2026 core update; personal finance -99%, fashion -95%. Building on a collapsing channel.
Micro-SaaS product	70% never pass $1k/mo; 40% never reach $1k MRR at all; median time to $1M ARR is 2 years 9 months. Wrong instrument for a 12-month revenue goal, and the cold-start distribution problem is my single greatest weakness.
An AI-visibility tracking tool	Already commoditized. $300M+ raised in the category since mid-2025; Profound at a $1B valuation on $155M; usable tools start at $25/mo. Never compete with venture-funded commodity infrastructure — **consume it instead**.
Shopify apps	17,600 apps, median under $1,000/mo, 5–12% *monthly* churn. Crowded and leaky.
ADA / accessibility remediation	Genuine regulatory demand, but DOJ pushed Title II to April 2027/2028, selling "compliance" carries direct liability, and DOJ explicitly cited generative AI's inability to automate remediation at scale. Keep as a Phase 3 adjacency, not the beachhead.
Upwork / Fiverr freelancing	Price-taking, platform-dependent, capped, and requires misrepresenting who does the work.

3. Customer and offer

3.1 Ideal customer profile

A B2B SaaS company that **already buys marketing** and is **demonstrably invisible in AI answers**:

- Series A–C, 20–200 employees, roughly $3M–$50M ARR
- Has a G2/Capterra presence and an active content program (proves budget and buy-in exist)
- Sells into a category with a real "best X software" query pattern
- Our scan shows <10% share of voice while a direct competitor exceeds 40%

Beachhead categories (chosen for vendor density, budget, low agency saturation, and — see §8 — insulation from the goods trade war): **legal tech** and **compliance/regtech**, with **cybersecurity** as the third. These serve buyers whose own demand is non-discretionary: regulatory deadlines and legal exposure do not soften when the economy does.

Revised August 2026. Earlier drafts named construction tech and field service as the beachhead. That is now the wrong call. Those vendors sell to contractors buying steel, copper, and building materials at 50% tariffs — margin compression flows straight through to software churn and frozen marketing budgets. HR/recruiting tech is demoted to secondary for a milder version of the same problem: it is levered to hiring, which is the first thing to slow. Expand once each category's research assets and case studies compound.

3.2 Offer ladder

Tier	Price	What it is	Purpose
AI Visibility Report	Free	Automated 120-query scan across 4 engines, share-of-voice vs 3 competitors, transcripts	The cold-email hook. Costs ~$0.40 to produce.
Diagnostic Sprint	**$1,500** one-time	2 weeks: 200-prompt tracked set, competitor citation teardown, source-gap map, 90-day roadmap	De-risks the buyer, produces cash immediately, converts to retainer at ~50%
Visibility Program	**$2,500/mo**	Tracked prompt set; 8–12 answer-engine content assets/mo; schema + entity hygiene; third-party profile coverage; monthly report	Core recurring product
Authority Program	**$5,000/mo**	Everything above + one original research study per quarter + digital PR pitching + comparison-page inclusion outreach	Attacks the 75% of citations that live off-site
Category Ownership	**$8,500/mo**	Multi-category/multi-geo, proprietary index or dataset, monthly study, quarterly strategy session	Ceiling-raiser for the best accounts

Contracts: month-to-month after a 3-month initial term, 30 days' notice. **Deliverable-based SLA — never an outcome guarantee.** We commit to activities and reporting, not to a citation rate we do not control.

3.3 Unit economics (Visibility Program, $2,500/mo)

Line	Monthly
Revenue	$2,500
LLM API — content production (Claude Sonnet 5 at $3/$15 per MTok, Opus 5 at $5/$25 for high-value assets)	~$110
Search / citation-tracking API (Perplexity Search API at $5 per 1,000 requests; ~960 tracked queries/mo)	~$10
Allocated shared tooling and infrastructure	~$30
Gross margin	**~$2,350 (94%)**

Customer acquisition cost is dominated by outbound infrastructure at roughly $150/month total, spread across every deal — well under $200 per client acquired at steady state. The dominant constraint is not money. It is delivery quality at volume, which Section 6 addresses.

---

4. Financial plan

Full month-by-month build, funnel assumptions, and scenario mechanics are in `FINANCIAL-MODEL.md`. Summary:

Scenario	New retainers/mo (steady)	Exit MRR	Year-1 revenue
Conservative	1.5	~$29,000	~$159,000
Base	**3.0**	**~$70,900**	**~$387,000**
Aggressive	4.0 @ higher tier mix (capacity-capped)	~$115,400	~$698,000

Base case assumes 3,000 emails/month at steady state, a 4.0% reply rate, 25% of replies positive, 22% of conversations closing a paid diagnostic, 50% diagnostic-to-retainer conversion, a $3,200 blended retainer, and **5% monthly churn**. It ends the year with roughly 22 active retainer clients.

Opening budget — the $1,000

Item	Cost
Entity formation	$0 — Manitoba corporation already held
US-domiciled USD account setup + cross-border accounting reserve	$150
Primary domain + 8 sending domains	$110
10 sending mailboxes, 3 months prepaid	$75
Smartlead, 3 months ($39/mo, unlimited mailboxes)	$117
LLM API credits	$200
Search / citation API credits	$50
Lead enrichment credits	$100
Brand assets, hosting, miscellaneous	$50
Unallocated reserve	$148
Total	**$1,000**

Website, tracking dashboard, outreach tooling, report generator, and every contract draft are built by me at no cash cost. Hosting is on a free tier. Stripe takes 2.9% + $0.30 with nothing upfront.

Reinvestment rule: through month 6, 100% of profit goes to sending capacity and API credits, in that order — outreach volume is the binding growth constraint. From month 7, cap outbound spend and route surplus into (a) original research studies that generate inbound and (b) a modest reserve against refunds. Take no distributions before month 9.

---

5. Go-to-market

Primary channel — report-led cold outbound. Build a target list of ~4,000 qualified companies by scanning public sources (G2 category pages, funding announcements, job postings for marketing roles). Score each one by running the free scan. Contact only companies whose scan shows a real, specific gap — this doubles as list hygiene and keeps spam complaints structurally low, because every email is genuinely relevant to its recipient.

Volume ramps as deliverability allows: 200/day maximum per new domain, 30-day ramp, targeting ~150/day steady state across ~10 mailboxes.

Secondary — publish the aggregate data. Every scan we run feeds a proprietary dataset: who gets cited, for what queries, on which engine, in which category. Quarterly we publish *The B2B AI Citation Index* from it. This does three things at once: it earns the citations and mentions that prove our own method works, it generates inbound, and it is the single most defensible asset the business builds. **The marketing channel is a live demonstration of the product.**

Tertiary — proof assets. Three free pilot clients in months 1–2, chosen for logo quality, in exchange for a case study and a testimonial. This is the fastest route past the no-track-record problem, and it is honest: real work, real results, disclosed relationship.

---

6. Operating model — how this actually runs

I do not run continuously. I run when invoked. Every recurring process must therefore be **code, not a Claude session**:

- Nightly scheduled jobs (GitHub Actions, free tier) run citation tracking for every client and every prospect, write results to a database, and flag movement.
- Report generation is templated and automated; I write the analysis and recommendations, not the assembly.
- The outreach queue is populated automatically; I write the copy.

Who	Does what	Time
Claude	All research, content, code, audits, reports, outreach copy, contract drafts, client email drafts	Invoked in batches
Principal (human)	Signs contracts, holds bank/Stripe/domain accounts, approves outbound before send, takes occasional calls, runs delivery sessions, final quality gate	3–6 hrs/week

The binding constraint is delivery quality at scale, not lead generation. At 22 clients the program produces 200–260 content assets per month. The correct response to demand beyond that is **to raise prices, not to add clients** — a 20% price increase and a 10% client reduction is strictly better for both margin and quality. Plan to raise prices at month 6 and month 10.

---

7. Legal, compliance, and ethics

Detail and the documents for counsel are in `../03-Legal-For-Review/3-Compliance-and-Ethics-Policy.md`. The essentials:

Entity — a Manitoba share corporation. The principal already holds one, and it is the right vehicle, not merely an acceptable one. Selling services into the US from a Canadian company is routine and does not meaningfully affect market access at this deal size. Three consequences matter:

- **Trading name: register, don't rename.** *Shortlist AIO* is registered as a **business name** of Glass House Gardens Inc. ($45 name reservation + $60 registration through the Manitoba Companies Office), not adopted by changing the corporate name via articles of amendment ($175). Manitoba's *Business Names Registration Act* makes the registration mandatory once we carry on business under a name other than the corporate one, which begins the moment the site goes live. Keeping the corporate name intact avoids re-papering the bank, CRA account, and any existing relationships; leaves the door open to hold several trade names under one corporation; and — given the plan's own month-4 kill criteria — avoids welding the legal entity to a brand that may not survive its first quarter. A trade name is a $105 reversible decision; a corporate name is not.

Note what this registration is *not*: it is a Manitoba compliance filing, **not brand protection**. It confers no trademark rights and does not stop anyone in the United States — where every customer is — from using the same name. If the brand proves worth defending, that is a separate US trademark question for a later year, not a launch task.

- **Trade war exposure: none, and the currency helps.** Tariffs are an instrument for goods; they do not apply to cross-border services or intangibles, and the September 2026 Canadian retaliation targets steel, dairy, agricultural equipment, appliances, paper, and electronics. Nothing we sell appears on any tariff schedule. Meanwhile a weak loonie (USD/CAD ~1.40) is a tailwind for a Canadian company invoicing in USD. Canada's Digital Services Tax — the one measure that would have reached a firm like this — was rescinded in 2025 and repealed retroactively, with refunds issued.
- **Tax.** Manitoba levies **0% provincial tax** on the first $500,000 of active business income for a Canadian-Controlled Private Corporation; with the 9% federal small-business rate that is a **9% combined rate**, the lowest in Canada alongside Yukon. The base case's ~$353,000 pre-tax profit sits entirely inside that band. This is a deferral advantage rather than an absolute one — personal tax applies when profits are distributed — but the plan reinvests through month 6 and holds reserves after, so the deferral is real money. *Condition:* CCPC status requires Canadian control. Without it the general rate applies (12% Manitoba + 15% federal = 27%).
- **No US federal income tax, no US withholding.** Under Articles V and VII of the Canada–US treaty, business profits are taxable in the US only through a US permanent establishment. Services performed from Canada for US clients create none. Separately, services performed outside the US are foreign-source income and are not subject to the 30% withholding — clients receive a **Form W-8BEN-E** and withhold nothing. The treaty position is not automatic: it requires a **timely Form 1120-F protective return with Form 8833 attached**, filed annually.
- **GST/HST works in our favour.** Advisory, consulting, and professional services supplied to non-resident, non-registered persons are **zero-rated** (Schedule VI, Part V, s. 23). We charge US clients 0%, and because zero-rated is not the same as exempt, we still claim input tax credits on Canadian expenses. Registration is

required above $30,000 CAD in four consecutive quarters and is worth doing voluntarily before then to recover ITCs.

Detail, and the questions for counsel and an accountant, are in `../03-Legal-For-Review/3-Compliance-and-Ethics-Policy.md`.

Hard ethical lines — non-negotiable, and also the correct business decision. The FTC's Rule on Consumer Reviews and Testimonials (effective 21 October 2024) carries civil penalties up to roughly **$52,000 per violation** and reaches conduct a company "knew or should have known" about. Accordingly:

- We never write, buy, sell, or incentivize reviews or testimonials — including AI-generated ones.
- We never post as a client or operate sockpuppet accounts. Community participation is done by named client employees with disclosed affiliation.
- We never build company-controlled sites presenting themselves as independent comparison or review sources.
- We disclose to every client, in the MSA, that delivery is AI-assisted.

This is not merely compliance posture. Astroturfing is the obvious shortcut in this category, it is the reason the category will eventually attract enforcement, and a firm that publishes its methodology and refuses to fake anything is differentiated on exactly the axis clients will start caring about.

Outbound. CAN-SPAM: accurate headers, honest subject lines, physical postal address, opt-outs honored within 10 business days. GDPR: B2B legitimate interest with a documented Legitimate Interest Assessment, disclosed data source, one-click opt-out. Sender rules: SPF, DKIM, and aligned DMARC; spam complaints held under 0.10%; 14-day minimum domain age; 30-day volume ramp.

Measurement. Citation tracking uses official paid APIs, never scraping against platform terms.

---

8. Risks and kill criteria

Risk	Severity	Mitigation
Cold outbound underperforms without a track record	**High**	Three free pilots for case studies in months 1–2; the report hook is the strongest available differentiator; kill criteria below
GEO commoditizes; tools automate the on-site layer	High	Move up the ladder into off-site earned mentions, original research, and PR — the 75% that cannot be automated by a dashboard
LLM platforms change citation mechanics	Medium	Contracts commit to activities, not outcomes; the tracking dataset detects shifts before clients do
Delivery quality degrades past ~20 clients	Medium	Cap client count; raise prices at months 6 and 10; automate assembly, never analysis

Risk	Severity	Mitigation
Client churn after 90 days without visible lift	Medium	Set expectations at 90 days in writing; report leading indicators (mention volume, source coverage) monthly, not just citations
Deliverability collapse	Medium	Ramp discipline, scan-qualified lists only, complaint monitoring, spare domains warmed in reserve
US–Canada trade war depresses SaaS marketing budgets	Medium	No direct exposure — tariffs apply to goods, not services — but a slowing economy is a real threat to the base case. Partly hedged by the nature of the pitch: we sell *reallocation* of an existing budget that is visibly underperforming, not a new line item. Reallocation pitches hold up better in a downturn than incremental-spend pitches. Beachhead verticals chosen for goods-exposure insulation (§3.1).
US buyer reluctance toward a Canadian supplier	Low	Evidence says B2B relationships are holding; a service carries no tariff cost to the buyer, so our nationality is commercially irrelevant to them. Neutral `.com` domain, USD pricing, and outreach that leads with their data rather than our location. Disclosed plainly in the footer, MSA, and privacy policy — not led with.
Non-tariff measures restrict access to US AI APIs	Low	Tail risk only: services cannot be tariffed, and negotiations are reportedly moving toward *more* cross-border data freedom. The engine layer is already pluggable adapters — losing one provider costs an engine, not the business.
Single-operator concentration	Low	Everything in version control and documented; no undocumented process

Kill criteria — decide, do not drift. If at the end of **month 4** the business has fewer than 2 paying retainers and under $8,000 cumulative revenue, stop selling and reassess. The most likely correction is a vertical change, not a business-model change: the same machine repoints at a different category in about two weeks.

---

9. What happens after year one

The two-year path is not "more retainers." It is:

1. **Productize the diagnostic** into a self-serve $299 report — same automated pipeline, no human in the loop, pure margin, and a top-of-funnel machine.
2. **Publish the Citation Index** as a durable, cited data asset that makes the firm the source rather than the vendor.
3. **Move up-market.** The same program sold to a $100M+ ARR company is a $15,000–$25,000/month engagement, and there are fewer than 30 firms in the world who will have a two-year citation dataset to sell alongside it.

Sources

Market and benchmark figures cited above are drawn from the research recorded in `RESEARCH-NOTES.md`, which lists each source URL against the claim it supports.

Shortlist AIO — Financial Model

Companion to `BUSINESS-PLAN.md`. All figures are modelled, not observed. The purpose of this document is to make every assumption explicit enough to be attacked.

---

1. Funnel assumptions (base case)

Stage	Rate	Source of the estimate
Emails sent / month at steady state	3,000	~150/day across 10 mailboxes; within Google/Microsoft ramp limits
Reply rate	**4.0%**	Published 2026 benchmarks are 3.1–3.7% on managed infrastructure; we model a modest premium for a bespoke, evidence-carrying email, not the 8–12% "elite" figure
Positive replies (of all replies)	25%	Conservative; most replies to any cold email are declines
Reply → real conversation	60%	
Conversation → paid Diagnostic ($1,500)	22%	
Diagnostic → retainer, within 45 days	50%	
Direct-to-retainer closes (no diagnostic)	~1/month	Larger prospects that skip the trial step
Net new retainers / month at steady state	**3.0**	
Blended retainer value	$3,200 /mo	Mix weighted toward the $2,500 tier
Monthly logo churn	**5.0%**	Services churn; results lag 60–90 days, so early churn is the real risk

At 3,000 emails → 120 replies → 30 positive → 18 conversations → 4 diagnostics → 2 retainer conversions + 1 direct = **3 net new retainers/month**.

---

2. Base case — month by month

MRR recursion: `MRR_end = MRR_prev × 0.95 + (new retainers × $3,200)`. Retainer revenue billed is taken as the mid-month average, which approximates start dates spread through the month.

Month	New retainers	Ending MRR	Active clients	Diagnostics sold	Retainer rev.	One-time rev.	Total revenue
1	0	$0	0	0	$0	$0	**$0**
2	0	$0	0	2	$0	$3,000	**$3,000**

Month	New retainers	Ending MRR	Active clients	Diagnostics sold	Retainer rev.	One-time rev.	Total revenue
3	1	$3,200	1	3	$1,600	$4,500	**$6,100**
4	2	$9,440	3	4	$6,320	$6,000	**$12,320**
5	3	$18,568	6	4	$14,004	$6,000	**$20,004**
6	3	$27,240	9	4	$22,904	$6,000	**$28,904**
7	3	$35,478	11	4	$31,359	$6,000	**$37,359**
8	3	$43,304	14	4	$39,391	$6,000	**$45,391**
9	3	$50,739	16	4	$47,022	$6,000	**$53,022**
10	3	$57,802	18	4	$54,271	$6,000	**$60,271**
11	3	$64,512	20	4	$61,157	$6,000	**$67,157**
12	3	**$70,886**	**22**	4	$67,699	$6,000	**$73,699**
				41	**$345,727**	**$61,500**	**$407,227**

Less a 5% allowance for refunds, prorations, and negotiated discounts:

> *### Year-1 recognized revenue: **~$387,000** ### Exit MRR: **$70,886** → **~$851,000 ARR run-rate***

---

3. Base case — cost side

Line	Basis
Outbound infrastructure	~$124/mo — 10 mailboxes ($25), Smartlead ($39), domains amortized ($10), lead enrichment ($50)
Prospect scanning	$200–300/mo — ~500 qualifying scans at ~$0.40 each
Client delivery (LLM + search APIs)	~$120/client/mo
Client-facing dashboard subscription	$160/mo from month 4, $320/mo from month 8 (we buy the commodity tool rather than build it)
Accounting / admin	$100/mo from month 6
Entity formation and registered agent	$150, month 1

	Amount
Total operating cost, year 1	~$21,800
Stripe processing (2.9% + $0.30)	~$11,800
Stripe FX conversion — **$0 if settling to a US-domiciled USD account**, ~$3,900 if not	$0

	Amount
Total cost	**~$33,600**
Pre-tax profit, year 1	**~$353,000**
Corporate tax — Manitoba CCPC, 9% combined on active business income under $500k	~$31,800
After-tax profit retained in the corporation	**~$321,000**

Currency. Revenue is invoiced in USD; the figures throughout this model are USD. At USD/CAD ~1.40 (August 2026), the base case's $387,000 is roughly **CAD $542,000**. Bank forecasts point to mild CAD appreciation toward 1.38 by Q4 2026, which is a modest headwind on translation and not material to any decision here. A trade-war-weakened loonie is, on balance, a tailwind for a Canadian company earning USD — and it raises the value of holding revenue in the US-domiciled USD account rather than converting on a schedule, so conversion timing stays a choice rather than an automatic haircut.

Revenue is invoiced in USD; the figures throughout this model are USD. Costs are also largely USD (APIs, sending tools, domains), so the currency exposure is limited to retained earnings. Settle Stripe into a US-domiciled USD account from the first invoice — otherwise automatic conversion takes roughly 1%, about $3,900 a year at base case.

The business is cash-positive from month 2 and never requires capital beyond the initial $1,000. This is the defining structural advantage: **there is no month in which growth is constrained by money.** It is constrained by delivery capacity and by close rate.

---

4. Scenarios

	Conservative	Base	Aggressive
Reply rate	3.0%	4.0%	5.0%
Net new retainers / mo (steady)	1.5	3.0	4.0
Blended retainer	$2,800	$3,200	$4,500
Monthly churn	6%	5%	4%
First retainer lands	Month 4	Month 3	Month 3
Diagnostics / mo	2	4	5
Exit MRR	**~$29,000**	**~$70,900**	**~$115,400**
Exit ARR run-rate	~$348,000	~$851,000	~$1,385,000
Year-1 revenue	**~$159,000**	**~$387,000**	**~$698,000**
Active clients at month 12	10	22	26

The aggressive case is **capacity-capped from month 10** (new adds cut from 4/mo to 2/mo). Uncapped it produces a higher number, but 30+ concurrent clients is not deliverable at the quality this business sells. Past that point the correct move is a price increase, not another logo.

---

5. Sensitivity — what actually moves the outcome

Year-1 revenue, varying one assumption at a time from the base case:

Change	Year-1 revenue	Δ
Reply rate 4.0% → 3.0%	~$300,000	–22%
Reply rate 4.0% → 5.0%	~$470,000	+21%
Churn 5% → 8%/mo	~$333,000	–14%
Churn 5% → 3%/mo	~$425,000	+10%
Blended retainer $3,200 → $4,000	~$450,000	+16%
First retainer slips month 3 → month 5	~$310,000	–20%

Two conclusions. First, *speed to first client matters as much as conversion rate* — a two-month slip costs as much as a full point of reply rate. Front-load the pilots. Second, **price is the cheapest lever available**: raising the blended retainer 25% costs nothing, consumes no capacity, and adds more than cutting churn by 40% does. Raise prices at months 6 and 10.

---

6. Reinvestment policy

Period	Rule
Months 1–6	100% of profit to sending capacity, then API credits. Outreach volume is the growth constraint.
Months 7–9	Cap outbound spend at ~$600/mo. Route surplus into original research studies (inbound engine) and a refund reserve equal to one month of MRR.
Months 10–12	Hold a 3-month operating reserve. Distributions permitted above that. Consider a part-time human editor if quality metrics slip.

A note on distributions. The 9% Manitoba CCPC rate is a *deferral*, not forgiveness — personal tax applies when profits leave the company as salary or dividends. Retaining earnings inside the corporation through year 1 is therefore worth real money, and the reinvestment schedule above already does that. Salary-versus-dividend mix is an accountant question to settle before the first distribution.

---

7. What would falsify this model

Watch these four numbers monthly. Each has a threshold that should trigger a change in plan, not a change in effort.

Metric	Healthy	Trigger
Reply rate	≥3%	<2% for two consecutive months → rewrite offer or change vertical
Spam complaint rate	<0.10%	>0.10% → stop sending, rebuild infrastructure
Diagnostic → retainer	≥40%	<25% → the diagnostic is not proving value; rebuild the deliverable
90-day client retention	≥80%	<60% → the program does not work; fix delivery before selling more

Kill criteria: fewer than 2 paying retainers and under $8,000 cumulative revenue at the end of month 4. Reassess the vertical first — the machine repoints in about two weeks.

Shortlist AIO — 90-Day Launch Plan

Everything below is either **[C]** — Claude does it, or **[H]** — the human principal must do it because it needs a legal person, a payment method, or a signature. The **[H]** column is the real schedule constraint; it totals roughly 4 hours a week.

Week 1 — Entity, identity, and the scanner

[H] (≈3 hrs, one time)

- Entity is already in place — Glass House Gardens Inc. No formation step.
- **Register** `Shortlist AIO` **as a business name** of the corporation through the

Manitoba Companies Office online portal — $45 name reservation + $60 registration. Required, not optional: Manitoba's *Business Names Registration Act* covers carrying on business under a name other than the corporate name, and the trade name hits the website, outbound email, and invoices from day one. Do **not** change the corporate name by articles of amendment ($175) — see `BUSINESS-PLAN.md` §7.

- Open a **US-domiciled USD account** (available to a Canadian-incorporated business on its existing registration documents) and connect it to Stripe as the settlement account. Doing this before the first invoice avoids ~1% on every dollar of revenue.
- Open the Stripe account in the corporate name; set USD as the presentment currency.
- **Set the Stripe statement descriptor to** `SHORTLISTAIO`**, not the corporate name.** A client who sees `GLASS HOUSE GARDENS` on a card statement for AI-visibility services will open a support ticket or, worse, file a chargeback. Invoices still carry the legal name.
- Book the cross-border accountant: GST/HST registration timing, and the Form 1120-F / 8833 protective-return calendar.
- Register the primary domain + 8 sending domains from the list I supply.
- Create the mailbox accounts and Smartlead account.

[C]

- Build the **citation scanner**: takes a company domain and a category, generates 40 buyer-intent prompts, queries ChatGPT, Perplexity, Gemini, and Google AI Overviews via official APIs, and scores share-of-voice against three named competitors. This is the single most important asset in the business — it is simultaneously the lead magnet, the sales weapon, the delivery instrument, and the reporting layer.
- Build the marketing site: positioning, offer ladder, three-page structure, one live sample report. Deploy to a free tier.
- Draft the MSA, SOW template, privacy policy, and terms for counsel.
- Configure SPF, DKIM, and aligned DMARC on every sending domain.

- Build the **CASL country screen** into the list pipeline: US-only for cold sequences, country verified from company records rather than inferred from the domain, with the exemption basis logged per recipient. This is a hard gate in code, not a checklist item.

Gate: scanner produces a correct, reproducible report for three test companies; Stripe settles a $1 test charge into the USD account with no conversion.

Week 2 — Proof before selling

[C]

- Build the target list: ~1,200 companies in the two beachhead categories, from public sources (G2 category pages, funding announcements, marketing-role job postings).
- Run scans across the list. Rank by *gap severity* — big gap, direct competitor dominant, evidence of existing marketing spend.
- Write the pilot outreach for the top 40.
- Start mailbox warmup. **No production sending until day 14 minimum** (Microsoft's domain-age floor), then ramp over 30 days.

[H] (≈1 hr)

- Send the pilot outreach to 40 hand-picked targets from the principal's own mailbox — not the cold infrastructure. Offer: a full diagnostic, free, in exchange for a case study and a testimonial if they like it.

Gate: 3 pilot clients accepted.

Weeks 3–4 — Deliver the pilots, build the machine

[C]

- Deliver all three pilot diagnostics in full. These set the quality bar for every paid engagement — over-invest here deliberately.
- Build the **delivery pipeline**: scheduled nightly tracking jobs, automated report assembly, content production templates, client dashboard.
- Write the case studies as the pilots produce results.
- Draft the first 6 outreach sequences (one per persona × category).

[H] (≈2 hrs)

- Review the MSA and the CASL screening procedure with counsel. Sign off on outbound copy.
- Sign the W-8BEN-E; keep a PDF ready to send the moment a US client's accounts-payable team asks.

- Approve the first sending batch.

Gate: pipeline runs unattended for 7 days without intervention.

Month 2 — Sell

[C]

- Scan-qualify and write outreach continuously. Ramp sending: 40/day week 5, 80/day week 6, 120/day week 7, 150/day week 8.
- Draft every reply, proposal, and SOW within an hour of a prospect responding — speed of response is a genuine competitive advantage here and costs nothing.
- Convert pilots to paid retainers at month-end.

[H] (≈4 hrs/week)

- Approve sends daily. Send proposals. Take calls where a prospect asks for one. Sign.

Targets: 880+ emails week 5–6, 2,000+ by week 8; **2 paid diagnostics closed**; 3 pilots converted or referenced.

Month 3 — Prove the loop repeats

[C]

- Deliver every diagnostic and retainer on schedule.
- Publish **The B2B AI Citation Index, Q4 2026** — the aggregate findings from every scan run to date. This is the inbound engine and it markets the firm through the exact mechanism the firm sells.
- Expand the target list to ~4,000. Open a third category if the first two are converting.
- Instrument the four health metrics from `FINANCIAL-MODEL.md` §7 into a weekly dashboard.

[H] (≈4 hrs/week)

- Sales calls, signatures, invoicing. First price review.

Targets: 3,000 emails/month steady state; **3 diagnostics + 1–2 retainers closed**; exit month 3 at ~$3,200 MRR and ~$6,100 in month-3 revenue.

Decision points

When	Question	Action if the answer is no
End of week 2	Did 3 pilots accept a *free* diagnostic?	The offer is not compelling. Rewrite it before spending on infrastructure — this is the cheapest possible failure.
End of month 2	Reply rate ≥3%? Complaints <0.10%?	Rewrite sequences; re-examine list quality before increasing volume.
End of month 3	≥1 paying retainer and ≥$8,000 cumulative?	Change vertical. Keep the machine.
End of month 4	**≥2 retainers and ≥$8,000 cumulative?**	**Kill criteria. Stop selling and reassess.**

Standing weekly rhythm after day 90

Day	Who	What
Monday	[C]	Weekly client tracking reports drafted; movement flagged
Monday	[H]	Approve and send; 30 min
Tue–Thu	[C]	Content production, outreach writing, prospect scans
Tue–Thu	[H]	Approve sends, handle replies, calls; ~2 hrs total
Friday	[C]	Health-metric dashboard, pipeline review, next week's queue
Friday	[H]	Review metrics, invoicing, banking; 45 min
Monthly	[C]	Client reports, one original research asset, price review input

Shortlist AIO — Research Notes & Sources

Every quantitative claim in the plan, with the source it came from. Read the caveat in §7 before treating any of these as precise.

1. B2B buying behavior has moved to AI assistants

Claim	Source
72% of B2B software buyers use ChatGPT in vendor evaluation; ChatGPT dominant at 63%	MarketScale
51% start research with an AI chatbot more often than Google; 71% use one somewhere in evaluation	G2 research via PR Newswire · Demand Gen Report
69% chose a different vendor on AI guidance; 33% bought from an unfamiliar vendor	MarketScale
Average shortlist compressed to 2.5 names from 3.2	AuthorityTech
51% of B2B tech brands have zero citations across ChatGPT, Perplexity, Gemini	Omnibound
94% of B2B buyers used AI tools in their most recent purchase	MR Research

2. The old channel is collapsing (why budget is available)

Claim	Source
58% average CTR decline for top organic results where an AI Overview appears (up from 34.5% in 2025)	SEO-Kreativ
60% of Google searches end without a click; AI Overviews on 48% of searches as of March 2026	Omnibound · Contently
71% of affiliate sites declined in the March 2026 core update; personal finance –99%, fashion –95%, health –93%	ClickDealer · Affiliyo

3. Market size and price points

Claim	Source
GEO services market ~$1.5–2.0B in 2026, $4B+ projected 2027; 45.5% CAGR to 2034	Intel Market Research · AuthorityTech
GEO retainers: SMB $1,500–5,000/mo; mid-market $2,000–8,000/mo; audits $500–2,000 one-time	Digital Elevator · RevvGrowth · WebFX

Claim	Source
Digital PR retainers $3,000–30,000/mo; single editorial placement $750–2,000; campaigns $3,000–20,000 for 10–40 mentions	Reporter Outreach · Loopex Digital
76% of SEOs pay $300+ per link; median secured price $150; willingness-to-pay averages $509	LinkBuild · OutreachDesk

4. What actually drives AI citations (the deliverable must work)

Claim	Source
Brand web mentions correlate **0.664** with AI citation rate vs backlinks at **0.218** (Ahrefs, 75,000 brands)	NetRanks
~75% of the citation equation lives on external domains	Contently
Heavy Reddit brand presence → 3.9× more ChatGPT citations (129K-domain study)	Otterly
46× difference in brand citation rate between platforms — ChatGPT 0.59%, Perplexity 13.05%; only 11% domain overlap	AuthorityTech · Otterly
Google AI Overviews show the strongest brand preference (59.8% of citations)	Leapd

5. Why the tooling layer was rejected, and the rejected alternatives

Claim	Source
AI visibility tools raised $300M+ mid-2025→spring-2026; Profound at $1B valuation on $155M raised; usable tools from $25/mo	Zapier · SearchIntel · Acromatico
Micro-SaaS: 70% under $1,000/mo; 40% never reach $1K MRR; median 2y9m to $1M ARR	Flowjam · VenturOS
Shopify: 17,600+ apps, median under $1,000/mo, 5–12% monthly churn	Week One Labs · Uptek
ADA Title II deadlines extended to April 2027/2028; DOJ cited genAI's inability to automate remediation at scale	Jackson Lewis · ADA Title III blog
AI automation agencies grew from ~2,000 (2024) to 12,000+ (2026) — saturation signal	Digital Applied

6. Distribution mechanics and compliance

Claim	Source
Cold email reply rates 3.1–3.7% on managed infrastructure; top 25% at 5.5%; volume blasts 0.45%	LaCleo · Amplemarket
Sender rules: SPF/DKIM/DMARC aligned, spam <0.10% (0.30% enforcement), 14-day domain age, 30-day ramp, 200/day cap on new domains	LeadHaste · Litemail
Lean cold email stack under $130/mo; Smartlead $39/mo with unlimited mailboxes	Litemail · Emelia
CAN-SPAM and GDPR legitimate-interest requirements for B2B prospecting; France Aug 2026 change is B2C only	Scrap.io · GrowLeads
FTC Rule on Consumer Reviews and Testimonials, effective 21 Oct 2024; ~$52,000/violation; "knew or should have known" standard	FTC · FTC Q&A · Wilson Sonsini
Entity formation: Wyoming $100 filing / $60 annual; Delaware $110 / $300 franchise tax; Stripe Atlas $500	BusinessAnywhere · SaveOffice

6A. Canadian entity — cross-border sources

Claim	Source
CASL is consent-based; penalties to CAD $10M per violation for corporations; director liability; private right of action; implied consent via conspicuously published address, role-relevant	McMillan LLP · CRTC FAQ · Smarte
Foreign-state exemption — ECPR para. 3(f): CEMs accessed in listed states incl. the US are exempt from CASL consent where they comply with that state's spam law	McInnes Cooper · Dorsey & Whitney
Manitoba small business rate 0%; 9% combined with federal on first $500k active business income — lowest in Canada with Yukon	Province of Manitoba · EY 2026 corporate rates · TaxTips
No US tax without a permanent establishment (Treaty Art. V/VII); requires timely Form 1120-F protective return + Form 8833	Edelkoort Smethurst CPAs · AKIF CPA
Services performed outside the US are foreign-source; no 30% withholding; W-8BEN-E is the documentation	Baker Tilly
Advisory/consulting/professional services to non-resident, non-registered persons are zero-rated (Sch. VI, Pt. V, s. 23); ITCs still claimable; $30k small-supplier threshold	CRA — Exports: Services and IPP · CRA Interpretation 222719
Stripe Canada charges USD but auto-converts to settlement currency at ~1%; a US-domiciled account avoids it; a USD account at a Canadian bank does not	Finofo · Loop

6B. Trade war exposure — sources (August 2026)

Claim	Source
50% US tariffs on ~$20bn of Canadian goods effective 22 Aug 2026, applying even to CUSMA-compliant goods, no expiry	Al Jazeera · Blakes timeline
Canada retaliating dollar-for-dollar from 8 Sep 2026, targeting steel, dairy, agricultural equipment, appliances, paper, electronics — **all goods**	CNBC · Al Jazeera
Tariffs do not apply to cross-border services or intangibles — the load-bearing fact for this business	Avalara · CRS, US–Canada Trade Relations
USD/CAD ~1.40–1.42 Aug 2026; five-bank consensus 1.38 by Q4	MTFX · LiteFinance
Canada's Digital Services Tax rescinded 2025, repealed retroactively via Bill C-15, CRA refunding ~$647M; Section 899 pressure was the cause	VATupdate · Canada.ca · ACTIA
"Digital trade alignment" in the negotiation package — likely *more* cross-border data freedom, not less	Michael Geist
B2B relationships largely holding; only ~12% of firms actively seeking alternative suppliers; Canadian Chamber reports no en-masse cold shoulder	CBC

7. Cost inputs for the unit economics

Input	Rate
Claude Sonnet 5	$3.00 / $15.00 per MTok in/out
Claude Opus 5	$5.00 / $25.00 per MTok in/out
Claude Haiku 4.5	$1.00 / $5.00 per MTok in/out
Perplexity Search API	$5 per 1,000 requests
Perplexity Sonar API	from $1/MTok, to $15/MTok output on Sonar Pro
OpenAI GPT-5.6 tiers	$5/$30 (Sol), $2/$12 (Terra), $0.20/$1.20 (Luna) per MTok

Sources: Anthropic pricing (current as of this plan); Perplexity API pricing; OpenAI API pricing.

---

8. Caveat on source quality — read this

A material share of the figures above come from **vendor and agency blogs**, which is what exists for a category this young. These sources have an interest in the numbers being large. Treat them as **order-of-magnitude indicators, not measurements**.

Three of them are load-bearing enough to verify independently before committing serious effort:

1. **The 51%-of-brands-have-zero-citations figure.** This is the premise of the entire cold-email hook. **It is also the one figure we can verify ourselves at near-zero cost** — the week-1 scanner measures it directly across the first 1,200 companies. If the real number is 15% rather than 51%, the addressable market shrinks by two-thirds and the offer needs rethinking. Run this check in week 2 and treat the result as authoritative over any published figure.
2. **The 4% reply rate.** Unverifiable until we send. This is why the plan front-loads three free pilots — they produce proof assets before the reply-rate assumption is tested at volume.
3. **GEO retainer price points.** Verifiable during sales conversations. If real mid-market willingness-to-pay is $1,200 rather than $2,500–8,000, the model's blended retainer is wrong and Year-1 revenue roughly halves.

The G2 and FTC sources are primary and reliable. The Ahrefs correlation study is reported secondhand here and is worth reading in the original before building the methodology on it.